
Image Classifier on the Fashion MNIST dataset

Abstract

This project shows an end-to-end implementation of image classification using Convolutional Neural Networks. The input data is a modified version of the Fashion-MNIST dataset where each image contains three articles and the sum of the prices is the label of the image. The VGG-16 structure was the most accurate amongst other types and was used to train a neural network. The *learning rate*, the *momentum* of the optimizer, the number of *epochs* and the *batch sizes* hyperparameters were tuned. The optimized values were $lr = 10^{-2}$, $momentum = 0.5$, $batch_size = 2^7$ and $n_epoch = 50$ which achieved 96.7% of accuracy on the validation set.

1 Introduction

In this report, an end-to-end project to implement and optimize a convolutional neural network is described. The input data used corresponds to a modified version of the Fashion MNIST dataset (1), on which each image analyzed contains three MNIST items from a group of five. The label of the observation is the sum of prices of the items on each image. To create a model that successfully predicts the targets of test images, a neural network with the architecture VGG-16 (2) is implemented and the original parameters are tuned to obtain the highest accuracy possible. The optimal parameters produced an accuracy of more than 97% on the validation set, thus providing a well-trained model for most practical tasks.

2 Dataset

The dataset used for training the proposed model corresponds to 60,000 images and their assigned labels. Each image contains three different articles of the Fashion MNIST original dataset (see some examples in Figure 1). These articles from the least to the most expensive are: T-shirt/top (\$1), Trouser (\$2), Pullover (\$3), Dress (\$4), Coat (\$5); where the number within parenthesis shows the price associated to each article. Each image is associated with a numerical label that is the sum of the articles' prices. Since each image contains articles from three different classes, the minimum value of the targets is \$5 and the maximum is \$13. A histogram with the nine different classes of the training set is shown in Figure 2. As illustrated, all classes are uniformly distributed and no re-sampling step is needed to analyze the data.



Figure 1: Examples of images on the input dataset

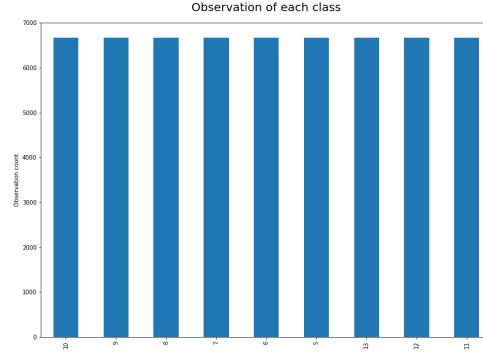


Figure 2: Histogram of training targets of the data.

3 Proposed approach

Each image contains a set of 128×64 gray-scale pixels in the (0, 255) range. The only pre-processing step conducted on the training dataset is re-scaling the range to (-1, 1). A fraction (20%) of the training dataset will be separated for testing, and the remaining data will be used for training the models.

The neural networks used in the model's development are based on the VGG-16 architecture (2) (discussed in subsection 3.1) and are implemented with *PyTorch* (3). The implementation of these neural networks is based on the tutorial (4) and (5).

3.1 Convolutional Network using VGG architecture

The architecture of VGG-16 is depicted in Figure 3 for an RGB figure of size 224×224 . The image is passed through a stack of convolutional layers using a filter of size 3×3 . The convolution stride is fixed to 1 pixel. Spatial pooling is carried out by five max-pooling layers, which follow some of the convolutional layers. Max pooling is performed over a 2×2 -pixel window with a stride of 2. Three fully connected layers follow the stack of convolutional layers. The final layer is a soft-max layer to provide results for the multi-class classification problem. All hidden layers are equipped with the RELU non-linearity. This convolutional network is renowned for its accuracy with image classification (6) and will be implemented for this project's task.




Figure 3: Architecture of VGG-16.

3.2 Models implementation

3.2.1 VGG-16 using *PyTorch*

The implementation of this neural network can be seen in Appendix A using *PyTorch*. It has been conducted such that the neural network architecture is easily adaptable for other versions of the VGG architecture (e.g. VGG-11), which accelerates the comparison between models.

3.2.2 Optimization algorithms - Stochastic Gradient Descent

For the neural network's training steps, *PyTorch*'s stochastic gradient descent (SGD) optimizer has been utilized. This well-known algorithm avoids local minima by randomly assigning initial weights for each training step. The SGD class receives as input the *learning rate* of the descent and its *momentum*, which can be used as hyperparameters for tuning the optimal model.

3.2.3 Hyperparameter tuning

The following parameters are tuned to find the optimal combination with the highest accuracy (smallest loss):

Learning rate: The learning rate of the SGD optimization algorithm, which controls how much to change the model in response to the estimated error each time the model weights are updated, is crucial. For example, a too-small value might lead to a slow convergence on loss, while a too big rate could yield an unstable training procedure. Thus, the learning rate is tested with $lr \in \{10^{-1}, 10^{-2}, 10^{-3}\}$ to confirm this hypothesis.

Momentum: The momentum of the SGD optimization algorithm, which accelerates gradients vectors, leads to faster converging. A system with momentum is proved to perform better in (7). Its value is set to *momentum* = 0.5.

Batch size: As SGD is sequential and uses small batches, the training time is inversely proportional to batch size in general. For a system with reasonable training time, a large batch size may be favourable. The *batch_size* $\in \{2^5, 2^7, 2^9\}$ were compared to confirm this hypothesis.

Number of epochs: In general, as the number of epochs increases, the model goes from underfitting, then to optimal state, then overfitting. Therefore, identifying a fair number of epochs is essential. The number of epochs in the range $n_epoch \in \{40, 60\}$ were compared.

As a baseline, the VGG-16 architecture an a SGD optimizer with the following hyperparameters value was used ($lr = 10^{-2}$, *momentum* = 0.5, *batch_size* = 2^7 , *n_epoch* = 60). From this base case, each parameter was changed one at a time and the loss and accuracy curves were analyzed for both datasets. Only a portion of these results is presented to highlight these experiments' main features due to space limits.

4 Results

The results discussed in this report are only a fraction of what is shown in the attached *mini_project_3_team_20*.

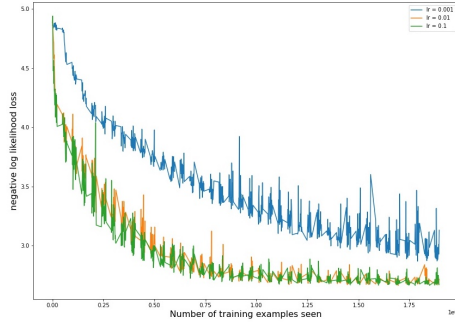
4.1 Analysis of hyperparameter tuning

4.1.1 Accuracy versus different learning rates

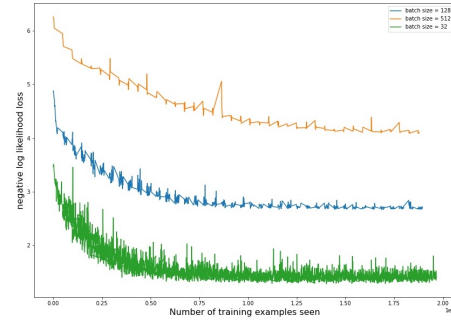
The values for the optimizer learning rate $lr \in \{10^{-1}, 10^{-2}, 10^{-3}\}$ are used for the same system to evaluate their impact. The results can be observed in Figure 4a. It can be seen that the loss curve for $lr = 10^{-3}$ converges very slowly, and after 40 epochs are still higher than the other two values. In addition, the loss curve of $lr = 10^{-1}$ is more fluctuating when compared to the curve of $lr = 10^{-2}$. Thus, the results agree with our previous hypothesis, and it can be concluded that the optimal learning rate among the tested values is $lr = 10^{-2}$.

4.1.2 Accuracy versus different batch sizes and time

The values *batch_size* $\in \{2^5, 2^7, 2^9\}$ are used on the same system to observe their effect, shown in Figure 4b. It can be observed that the smaller the batch size, the faster the convergence of the loss curve. Besides, the training versus batch size is given in Table 1. The training time is not proportional to batch size, and *batch_size* = 2^7 yields the best training time. As a model with good accuracy and reasonable training time is wanted, it can be concluded that the optimal batch size among the tested values is *batch_size* = 2^7 .



(a) Versus different learning rates

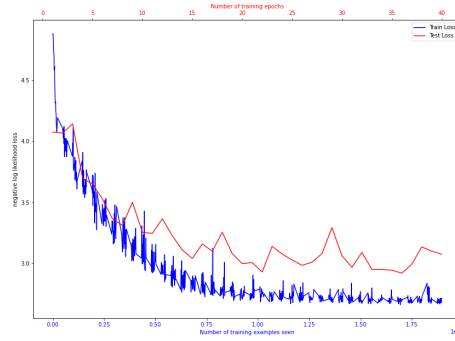


(b) Versus different batch sizes


Figure 4: Train loss variation

Table 1: Average time per epoch of different batch size

Batch size (samples)	32	128	512
Total time for 40 epochs (secs)	6402.2	5603.5	6182.6
Total time for 60 epochs (secs)	9605.1	8401.0	9278.9
Average time per epoch (secs)	160.0	139.9	154.5



(a) $n_epoch = 40$



(b) $n_epoch = 60$

Figure 5: Train and test losses with n_epoch

4.1.3 Accuracy versus different numbers of epochs

The previous analyses showed the behaviour of the *training* loss curve with different hyperparameters. Yet, there is no use in comparing only training curves when the model predicts unknown values. To select optimal models, curves of losses and accuracies on a test set should be explored. Figure 5 shows these accuracies curves for the cases $n_epoch \in \{40, 60\}$. In fact, both training loss curves fluctuation decrease after 30 epochs; however, the test loss curves remain steady only after 50 epochs. Therefore, to minimize the risk of over-fitting, the optimal value for the number of epochs should be $n_epoch = 50$.

4.2 Analysis of the final result

Based on previous observations and analysis, it can be concluded that the optimal hyperparameters for the VGG-16 network is $lr = 10^{-2}$, $momentum = 0.5$, $batch_size = 2^7$ and $n_epoch = 50$. A new CNN is trained with these hyperparameters, and the corresponding accuracy is illustrated in Figure 6, which achieved 96.7% over a test set of 12,000 images. Besides, from Figure 7, it can be observed that most error classes are mislabeled only to the nearest class, which suggests that our model only mistakes between very similar images.

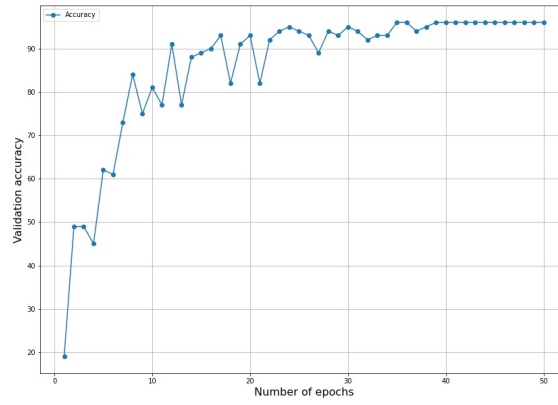


Figure 6: Accuracy over iteration obtained with $lr = 10^{-2}$, $momentum = 0.5$, $batch_size = 2^7$ and $n_epoch = 50$.

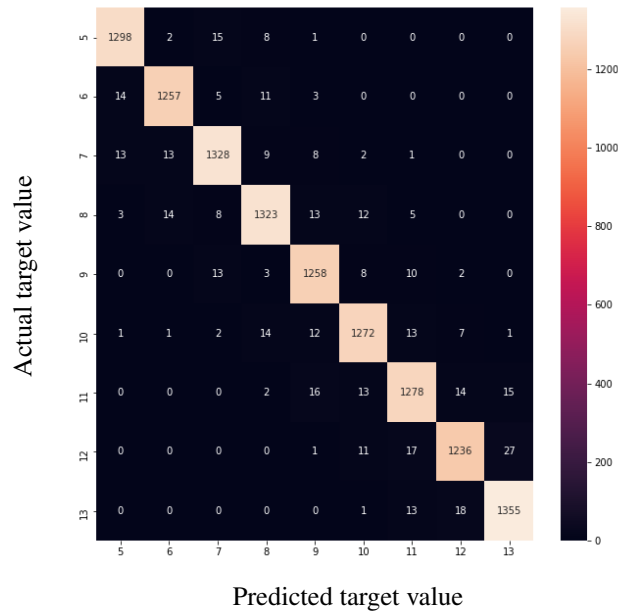


Figure 7: Heat Map of the Confusion Matrix for results obtained with $lr = 10^{-2}$, $momentum = 0.5$, $batch_size = 2^7$ and $n_epoch = 50$.

5 Conclusion

This project presents the results of applying a dense neural network of architecture VGG-16 to the task of image recognition and classification. The pre-processing of the images and the analysis of their distribution was shown. The composition of this popular neural network was briefly discussed. Several figures of interest are analyzed throughout the project: running time of the network, loss and accuracy vs observed data for both test and training set. Different hyperparameters for the model influenced these figures. These experiments and comparisons were used to select an optimal model with the highest accuracy possible on the validation set.

References

- [1] H. Xiao, K. Rasul, and R. Vollgraf, “Fashion-mnist: A novel image dataset for benchmarking machine learning algorithms,” *arXiv preprint arXiv:1708.07747*, 2017.
- [2] K. Simonyan and A. Zisserman, “Very deep convolutional networks for large-scale image recognition,” *arXiv preprint arXiv:1409.1556*, 2014.
- [3] FaceBook, “*PyTorch*,” 2020. <https://pytorch.org/> [Accessed: 11/28/2020].
- [4] A. Persson, “*Pytorch VGG implementation from scratch*,” 2020. <https://www.youtube.com/watch?v=ACmuBbuXn20> [Accessed: 11/29/2020].
- [5] A. Kak and C. Bouman, “*Lecture Notes on Deep Learning*,” February 2020. <https://engineering.purdue.edu/DeepLearn/pdf-kak/week5.pdf> [Accessed: 11/30/2020].
- [6] J. Wei, “Vgg neural networks: The next step after alexnet,” July 2019. <https://towardsdatascience.com/vgg-neural-networks-the-next-step-after-alexnet-3f91fa9ffe2c> [Accessed: 11/29/2020].
- [7] G. D. Ilya Sutskever¹, James Martens and G. Hinton, “On the importance of initialization and momentum in deep learning,” *Proceedings of the 30th International Conference on Machine Learning*, pp. 1139–1147, 2013.

A Appendix

Complete notebook with figures can be found in the attached files MiniProject3Team20.ipynb.

It follows the code of implementing a VGG-16 network to this project.

```
1 # -*- coding: utf-8 -*-
2 """Mason_Mini-Project_3_group_20.ipynb
3
4 Automatically generated by Colaboratory.
5
6 Original file is located at
7     https://colab.research.google.com/drive/1
8     P6Xhs3G0G0Rq8lDVG96oweIX8X94M
9
10 # Introduction
11 In the following you will see how to read the provided files for the
12 mini-project 3.
13 First you will see how to read each of the provided files. Then, you
14 will see a more elegant way of using this data for training neural
15 networks.
16 """
17
18 # Commented out IPython magic to ensure Python compatibility.
19 from google.colab import drive
20 from google.colab import files
21
22 drive.mount('/content/gdrive')
23 # %cd '/content/gdrive/MyDrive/Colab Notebooks/Mini Project 3 ECSE 551
24     Team 20'
25
26 import pickle
27 import matplotlib.pyplot as plt
28 import numpy as np
29 from torchvision import transforms
30 from torch.utils.data import Dataset
31 from torch.utils.data import DataLoader
32 from PIL import Image
33 import torch
34
35 path_lib = '/content/gdrive/MyDrive/Colab Notebooks/Mini Project 3
36     ECSE 551 Team 20/'
37 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
38 #device = 'cpu'
39
40 """
41 Let's see how the data looks like.
42 """
43
44 import pandas as pd
45 import matplotlib.pyplot as plt
46
47 train = pd.read_csv("./data/TrainLabels.csv")
48
49 plt.subplots(figsize=(14, 10))
50 train['class'].value_counts().plot.bar()
51 plt.ylabel("Observation count")
52 plt.title('Observation of each class', fontsize=20, pad=20)
53 plt.savefig(path_lib + 'graph/class_hist.png')
54 plt.show()
55
56 print((train['class'].value_counts()))
57
58 print("Each class has the same number of observation +-1 data")
59
60 # Read a pickle file and display its samples
```

```

55 # Note that image data are stored as unit8 so each element is an
    integer value between 0 and 255
56 data = pickle.load( open( './data/Train.pkl', 'rb' ), encoding='bytes'
    )
57 targets = np.genfromtxt('./data/TrainLabels.csv', delimiter=',',
    skip_header=1)[: ,1:]
58 # plt.subplots(figsize=(14, 10))
59 plt.imshow(data[1234,:,:], cmap='gray', vmin=0, vmax=256)
60 print(data.shape, targets.shape)
61
62 """# Dataset class
63 *Dataset* class and the *Dataloader* class in pytorch help us to feed
    our own training data into the network. Dataset class is used to
    provide an interface for accessing all the training or testing
    samples in your dataset. For your convinance, we provide you with
    a custom Dataset that reads the provided data including images (.
    pkl file) and labels (.csv file).
64
65 # Dataloader class
66 Although we can access all the training data using the Dataset class,
    for neural networks, we would need batching, shuffling,
    multiprocessing data loading, etc. DataLoader class helps us to do
    this. The DataLoader class accepts a dataset and other parameters
    such as batch_size.
67 """
68
69 training_phase = True
70 export = False
71 idx_input = None
72 re_train = True
73
74 """
75 If in training phase, set the following parameters
76 """
77
78 train_size = 48000
79 test_size = 12000
80 batch_size = 2 ** 7 #feel free to change it
81
82 #Optimizer calibratable hyper parameters
83 c_lr = 0.001
84 c_momentum = 0.5
85
86 #Network iteration calibratable paramerters
87 n_epochs = 40
88
89 # Transforms are common image transformations. They can be chained
    together using Compose.
90 # Here we normalize images img=(img-0.5)/0.5
91 img_transform = transforms.Compose([
92     transforms.ToTensor(),
93     transforms.Normalize([0.5], [0.5])
94 ])
95
96 # img_file: the pickle file containing the images
97 # label_file: the .csv file containing the labels
98 # transform: We use it for normalizing images (see above)
99 # idx: This is a binary vector that is useful for creating training
    and validation set.
100 # It return only samples where idx is True
101
102 class MyDataset(Dataset):
103     def __init__(self, img_file, targets, transform=None, idx = None):
104         self.data = pickle.load( open( img_file, 'rb' ), encoding='
            bytes')
105         self.targets = targets

```



```

106         # self.targets = np.genfromtxt(label_file, delimiter=',',
skip_header=1)[: ,1:]
107         if idx is not None:
108             self.targets = self.targets[idx]
109             self.data = self.data[idx]
110         if transform is not None:
111             self.transform = transform
112
113     def __len__(self):
114         return len(self.targets)
115
116     def __getitem__(self, index):
117         img, target = self.data[index], int(self.targets[index])
118         img = Image.fromarray(img.astype('uint8'), mode='L')
119
120         if self.transform is not None:
121             img = self.transform(img)
122
123         return img, target
124
125 class MyValidationSet(Dataset):
126     def __init__(self, img_file, transform=None, idx = None):
127         self.data = pickle.load( open( img_file, 'rb' ), encoding='
bytes')
128         # self.targets = np.genfromtxt(label_file, delimiter=',',
skip_header=1)[: ,1:]
129         if idx is not None:
130             self.data = self.data[idx]
131         if transform is not None:
132             self.transform = transform
133
134     def __len__(self):
135         return len(self.data)
136
137     def __getitem__(self, index):
138         img = self.data[index]
139         img = Image.fromarray(img.astype('uint8'), mode='L')
140
141         if self.transform is not None:
142             img = self.transform(img)
143
144         return img
145
146 # Read image data and their label into a Dataset class
147
148 if training_phase:
149     print("Training phase dataset generation")
150     # Split the data into the training set and the test set
151     from sklearn.model_selection import ShuffleSplit
152     SS = ShuffleSplit(n_splits=1, train_size=train_size, test_size=
test_size, random_state=28)
153     targets = np.genfromtxt('./data/TrainLabels.csv', delimiter=',',
skip_header=1)[: ,1:]
154     train_index, test_index = next(SS.split(targets))
155
156     #Create the train dataset
157     target_csv = np.genfromtxt('./data/TrainLabels.csv', delimiter=',',
skip_header=1)[: ,1:]
158     min_target = min(target_csv)
159     target_csv = target_csv - min_target
160
161     train_dataset = MyDataset('./data/Train.pkl', target_csv, idx=
train_index, transform=img_transform)
162     test_dataset = MyDataset('./data/Train.pkl', target_csv, idx=
test_index, transform=img_transform)
163

```

```

164 print(test_dataset.targets)
165 #test_dataset = MyDataset('./Train.pkl', './TrainLabels.csv', idx=
    test_index, transform=img_transform)
166
167 #Apply the dataloader
168 train_loader = DataLoader(train_dataset ,batch_size=batch_size,
    shuffle=True)
169 test_loader = DataLoader(test_dataset ,batch_size=batch_size,
    shuffle=False)
170
171 examples = enumerate(test_loader)
172 batch_idx, (example_data, example_targets) = next(examples)
173 example_data = np.squeeze(example_data)
174
175 fig = plt.figure()
176 for i in range(6):
177     plt.subplot(2,3,i+1)
178     #plt.tight_layout()
179     plt.imshow(example_data[i].cpu().numpy(), cmap='gray', vmin=-1,
    vmax=1, interpolation='none')
180     plt.title("Ground Truth: {}".format(example_targets[i]))
181     plt.xticks([])
182     plt.yticks([])
183 fig
184 print(example_data.shape)
185
186 else:
187     print("Not in training phase")
188     #Create the train dataset
189     target_csv = np.genfromtxt('./data/TrainLabels.csv', delimiter=',',
    skip_header=1)[: ,1:]
190     min_target = min(target_csv)
191     target_csv = target_csv - min_target
192
193     train_dataset = MyDataset('./data/Train.pkl', target_csv, idx=None,
    transform=img_transform)
194     train_loader = DataLoader(train_dataset, batch_size=batch_size,
    shuffle=True)
195
196     #Create the test dataset
197     test_dataset = MyValidationSet('./data/Test.pkl', idx=None,
    transform=img_transform)
198     test_loader = DataLoader(test_dataset ,batch_size=batch_size,
    shuffle=True)
199
200 """# Define our NN
201 Define the VGG16 to be used in this project
202 """
203
204 import torch.nn as nn
205 import torch.nn.functional as F
206 import torch.optim as optim
207
208 class Net(nn.Module):
209     # This part defines the layers
210     def __init__(self):
211         super(Net, self).__init__()
212         # At first there is only 1 channel (greyscale). The next
    channel size will be 10.
213         input_sz_h = 128
214         input_sz_v = 64
215         fltr_sz_cv_1 = 3
216         fltr_num_cv_1 = 10
217         fltr_sz_cv_2 = 3
218         fltr_num_cv_2 = 10

```

```

219         final_sz_h = np.floor(((input_sz_h-fltr_sz_cv_1+1)/2-
220         fltr_sz_cv_2+1)/2)
221         final_sz_v = np.floor(((input_sz_v-fltr_sz_cv_1+1)/2-
222         fltr_sz_cv_2+1)/2)
223         self.img_sz = int(fltr_num_cv_2 * final_sz_h * final_sz_v)
224
225         self.conv1 = nn.Conv2d(1, fltr_num_cv_1, kernel_size=
226         fltr_sz_cv_1)
227         self.conv2 = nn.Conv2d(fltr_num_cv_1, fltr_num_cv_2,
228         kernel_size=fltr_sz_cv_2)
229         self.conv2_drop = nn.Dropout2d()
230
231         NN_neurons = 120
232         self.fc1 = nn.Linear(self.img_sz, NN_neurons)
233         self.fc2 = nn.Linear(NN_neurons, 9)
234
235         # And this part defines the way they are connected to each other
236         # (In reality, it is our foreward pass)
237         def forward(self, x):
238
239             x = F.relu(F.max_pool2d(self.conv1(x), 2))
240             x = F.relu(F.max_pool2d(self.conv2(x), 2))
241             #x = x.view(-1, self.img_sz)
242             x = x.view(x.size(0), -1)
243             x = F.relu(self.fc1(x))
244             x = F.dropout(x, training=self.training)
245             x = self.fc2(x)
246
247             return F.log_softmax(x,dim=0)
248
249     class Net_VGG16(nn.Module):
250         # This part defines the layers
251         def __init__(self):
252             super(Net_VGG16, self).__init__()
253
254         # architecture based on https://engineering.purdue.edu/DeepLearn/pdf-
255         kak/week5.pdf
256         #VGG16 = {'model':[64, 64, 'M', 128, 128, 'M', 256, 256, 256, 'M',
257         512, 512, 512,'M', 512, 512, 512, 'M'], 'name':'VGG11'}
258
259         self.conv_seqn = nn.Sequential(
260             # Conv Layer block 1:
261             nn.Conv2d(in_channels=1, out_channels=64, kernel_size=3,
262             stride=1, padding=1),
263             nn.BatchNorm2d(64),
264             nn.ReLU(inplace=True),
265             nn.Conv2d(in_channels=64, out_channels=64, kernel_size=3,
266             stride=1, padding=1),
267             nn.BatchNorm2d(64),
268             nn.ReLU(inplace=True),
269             nn.MaxPool2d(kernel_size=2, stride=2),
270             # Conv Layer block 2:
271             nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3,
272             stride=1, padding=1),
273             nn.BatchNorm2d(128),
274             nn.ReLU(inplace=True),
275             nn.Conv2d(in_channels=128, out_channels=128, kernel_size=3,
276             stride=1, padding=1),
277             nn.BatchNorm2d(128),
278             nn.ReLU(inplace=True),
279             nn.MaxPool2d(kernel_size=2, stride=2),
280             # Conv Layer block 3:

```

```

275         nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3,
276                   stride=1, padding=1),
277         nn.BatchNorm2d(256),
278         nn.ReLU(inplace=True),
279         nn.Conv2d(in_channels=256, out_channels=256, kernel_size=3,
280                   stride=1, padding=1),
281         nn.BatchNorm2d(256),
282         nn.ReLU(inplace=True),
283         nn.MaxPool2d(kernel_size=2, stride=2),
284         # Conv Layer block 4:
285         nn.Conv2d(in_channels=256, out_channels=512, kernel_size=3,
286                   stride=1, padding=1),
287         nn.BatchNorm2d(512),
288         nn.ReLU(inplace=True),
289         nn.Conv2d(in_channels=512, out_channels=512, kernel_size=3,
290                   stride=1, padding=1),
291         nn.BatchNorm2d(512),
292         nn.ReLU(inplace=True),
293         nn.Conv2d(in_channels=512, out_channels=512, kernel_size=3,
294                   stride=1, padding=1),
295         nn.BatchNorm2d(512),
296         nn.ReLU(inplace=True),
297         nn.MaxPool2d(kernel_size=2, stride=2),
298         # Conv Layer block 5:
299         nn.Conv2d(in_channels=512, out_channels=512, kernel_size=3,
300                   stride=1, padding=1),
301         nn.BatchNorm2d(512),
302         nn.ReLU(inplace=True),
303         nn.Conv2d(in_channels=512, out_channels=512, kernel_size=3,
304                   stride=1, padding=1),
305         nn.BatchNorm2d(512),
306         nn.ReLU(inplace=True),
307         nn.MaxPool2d(kernel_size=2, stride=2),
308     )
309     self.to_linear = 4096
310     c_num_classes = 9
311
312     self.fc_seqn = nn.Sequential(
313         nn.Linear(self.to_linear, 4096),
314         nn.ReLU(inplace=True),
315         nn.Dropout(p=0.5),
316         nn.Linear(4096, 4096),
317         nn.ReLU(inplace=True),
318         nn.Dropout(p=0.5),
319         nn.Linear(4096, c_num_classes)
320     )
321
322     # And this part defines the way they are connected to each other
323     # (In reality, it is our foreward pass)
324     def forward(self, x):
325
326         x = self.conv_seqn(x)
327         x = x.view(x.size(0), -1)
328         x = self.fc_seqn(x)
329
330         return F.log_softmax(x, dim=0)

```

```

332
333 #network = Net().to(device)
334 network = Net_VGG16().to(device)
335
336 optimizer = optim.SGD(network.parameters(), lr=c_lr, momentum=
    c_momentum)
337
338 train_losses = []
339 train_counter = []
340 test_losses = []
341 test_counter = range(1,n_epochs+1)
342 predi = []
343
344 def train(epoch):
345     network.train()
346     for batch_idx, (data, target) in enumerate(train_loader):
347         target = target.to(device)
348         data = data.to(device)
349         optimizer.zero_grad()
350         output = network(data)
351         loss = F.nll_loss(output, target) #negative log liklhood loss
352         loss.backward()
353         optimizer.step()
354
355         if batch_idx % 20 == 0:
356             print('Train Epoch: {} [{}/{} ( {:.0f}% )]\tLoss: {:.6f}'.format(
357                 epoch, batch_idx * len(data), len(train_loader.dataset),
358                 100. * batch_idx / len(train_loader), loss.item()))
359             train_losses.append(loss.item())
360             train_counter.append(
361                 (batch_idx*64) + ((epoch-1)*len(train_loader.dataset)))
362
363     torch.save(network.state_dict(), './models/model.pth')
364     torch.save(optimizer.state_dict(), './models/optimizer.pth')
365
366 def test():
367     network.eval()
368     test_loss = 0
369     correct = 0
370     predi = []
371     with torch.no_grad():
372         for batch_idx, (data, target) in enumerate(test_loader):
373
374             target = target.to(device)
375             data = data.to(device)
376             output = network(data)
377             test_loss += F.nll_loss(output, target, reduction='sum').item()
378             pred = output.data.max(1, keepdim=True)[1]
379             correct += pred.eq(target.data.view_as(pred)).sum()
380             predi.append(pred)
381             #print(pred)
382             # print(len(predi[0]))
383             # print(len(predi)) # this should
384
385     test_loss /= len(test_loader.dataset)
386     test_losses.append(test_loss)
387     print('\nTest set: Avg. loss: {:.4f}, Accuracy: {}/{} ( {:.0f}% )\n'.
        format(
388         test_loss, correct, len(test_loader.dataset),
389         100. * correct / len(test_loader.dataset)))
390     return pred, predi
391
392 def test_val():
393     network.eval()
394     test_loss = 0
395     correct = 0

```

```

396 predi = []
397 with torch.no_grad():
398     for data in test_loader:
399
400         #target = target.to(device)
401         data = data.to(device)
402         output = network(data)
403         # test_loss += F.nll_loss(output, target, size_average=False).
item()
404         pred = output.data.max(1, keepdim=True)[1]
405         #correct += pred.eq(target.data.view_as(pred)).sum()
406         predi.append(pred)
407         #print(pred)
408 # print(len(predi[0]))
409 # print(len(predi)) # this should
410
411 # test_loss /= len(test_loader.dataset)
412 # test_losses.append(test_loss)
413 # print('\nTest set: Avg. loss: {:.4f}, Accuracy: {}/{} ({:.0f}%) \n
'.format(
414     # test_loss, correct, len(test_loader.dataset),
415     # 100. * correct / len(test_loader.dataset)))
416 return pred, predi
417
418 def load_network():
419     network.load_state_dict(torch.load('./models/model.pth'))
420     optimizer.load_state_dict(torch.load('./models/optimizer.pth'))
421
422 if training_phase:
423     if re_train:
424         for epoch in range(1, n_epochs+1):
425             train(epoch)
426             pred, predi = test()
427     else:
428         load_network()
429         pred, predi = test()
430 else:
431     if re_train:
432         for epoch in range(1, n_epochs+1):
433             train(epoch)
434             pred, predi = test_val()
435     else:
436         load_network()
437         pred, predi = test_val()
438
439 """# Visualize result"""
440
441 if re_train:
442     fig, ax1 = plt.subplots(figsize=(14, 10))
443
444     ax1.set_xlabel('Number of training examples seen', color='blue')
445     ax1.set_ylabel('negative log likelihood loss')
446     lns1 = ax1.plot(train_counter, train_losses, color='blue', label='
Train Loss')
447     ax1.tick_params(axis='x', labelcolor='blue')
448
449     plt.legend(['Train Loss', 'Test Loss'], loc='upper right')
450     plt.savefig('train_loss.png')
451
452 if training_phase:
453     ax2 = ax1.twinx()
454     lns2 = ax2.plot(test_counter, test_losses, color='red', label='
Test Loss')
455     ax2.set_xlabel('Number of training epochs', color='red')
456     ax2.tick_params(axis='x', labelcolor='red')

```

```

457     lns = lns1+lns2
458     labs = [l.get_label() for l in lns]
459
460     ax2.legend(lns, labs, loc='upper right')
461     plt.show
462
463     plt.savefig(path_lib + 'graph/train_test_loss.png')
464
465 file = open('./result/lr0001.pkl', 'wb')
466 pickle.dump([train_counter, train_losses, test_losses], file)
467 file.close()
468
469 file = open('./result/lr0001.pkl', 'rb')
470 train_counter2, train_losses2, test_losses2 = pickle.load(file)
471 plt.plot(train_counter2, train_losses2)
472
473 """
474 This creates a single list out of the Predi variable.
475 Since Predi is a nested list, each elements are appended to the new
476 list predictions
477 Since the batch size doesnt exactly count for the total number of
478 values,
479 this deletes all values above the size of the test_dataset
480 """
481 batch_size = len(predi[0])
482 test_size = len(test_dataset)
483
484 predictions = []
485 for list in predi:
486     for item in list:
487         predictions.append(item.item() + min_target)
488 #print(len(predictions))
489
490 del predictions[test_size:]
491 #print(len(predictions))
492 #print(predictions)
493
494 """
495 Create a Pandas Datafram to see the data
496 It is also easier to export
497 """
498 result=pd.DataFrame(predictions)
499 result.reset_index(level=0, inplace=True)
500 result.columns = ["id", "class"]
501
502 if not training_phase:
503     if re_train:
504         result.to_csv("./result/ECSE551_Group20.csv", index=False)
505     else:
506         result.to_csv("./result/ECSE551_Group20_.csv", index=False)
507
508 # files.download("ECSE551_Group20_.csv")
509
510 #print(predi[0][0].item())
511
512 fig = plt.figure()
513 result['class'].value_counts().plot.bar()
514 plt.ylabel("Observation count")
515 plt.show()
516
517 if training_phase:
518     # https://towardsdatascience.com/multi-class-text-classification-
519     # with-scikit-learn-12f1e60e0a9f
520     """

```

```

520 Let's build a confusion matrix
521 """
522 import seaborn as sns
523 from sklearn.metrics import confusion_matrix
524
525 # conf_mat = confusion_matrix(test_dataset.targets + min_target,
526                             # predictions)
527 conf_mat = confusion_matrix(test_loader.dataset.targets + min_target
528                             , predictions)
529 fig, ax = plt.subplots(figsize=(10,10))
530
531 unique_price = np.unique(test_loader.dataset.targets + min_target).
532                   astype('int32')
533 sns.heatmap(conf_mat, annot=True, fmt='d', xticklabels=unique_price,
534             yticklabels=unique_price)
535 plt.savefig(path_lib + 'graph/heatmap.png')
536
537 # sns.heatmap(conf_mat, annot=True, fmt='d',
538 #             xticklabels=test_dataset.targets, yticklabels=
539 #                 test_dataset.targets)
540
541 plt.ylabel('Actual')
542 plt.xlabel('Predicted')
543 plt.show()

```